

A Safeguarded Data-Analyst Agent for the Palmer Penguins Dataset

Agentic AI Systems Analysis Report

Dennis O'Higgins

October 2026

1. Overview

This project designs and implements a single semi-autonomous agent that answers plain-English questions about a real scientific dataset, the Palmer Penguins measurements of 344 Antarctic penguins (Gorman et al., 2014). Following the usual definition of an agent as a system that perceives its situation and acts on it with some autonomy (Wooldridge & Jennings, 1995), the agent receives a request, decides whether it is in scope, plans one to six analysis steps, calls approved tools, remembers recent turns, and returns an answer together with a trace of how it was reached. Recent surveys describe the same loop for language-model agents of profile, memory, planning and action (Wang et al., 2024). The main design decision is that the reasoning engine is an explicit rule-based planner written in Python, not a language model. I chose this so that every decision can be inspected and every result reproduced. An optional language-model planner is included behind the same interface, but it was **not executed** in this project because no credentials were available, so no claim is made about its behaviour. The report covers the use case, architecture, decision logic, tools and memory, an evaluation on 28 test tasks, ethical considerations, limitations and future work.

2. Task and Use Case Description

The use case is exploratory data analysis. A user wants to know, for example, which species is heaviest, whether flipper length relates to bill depth, or how many penguins live on each island, without writing code. The dataset has 344 rows and seven columns: species, island, four numeric measurements (bill length, bill depth, flipper length, body mass) and sex, with 2 rows missing all four measurements and 11 missing sex. It was downloaded from the seaborn-data repository, and the notebook verifies its SHA-256 checksum on every run. I did not verify the dataset's licence, so users should check the source before redistributing it.

The agent is **single-agent by design**. The workload is a short sequential chain over one dataset and one memory, so splitting it across several agents would add message passing and failure points without adding capability. It is **semi-autonomous**: it plans and executes multi-step analyses without prompting, but it asks the user rather than guessing when a request is ambiguous and refuses requests outside its scope. Its permitted actions are limited to six read-only tools.

3. Agent Architecture and Workflow Design

System architecture and design. Figure 1 shows the components. A request first passes a safety screen, then the planner, then a plan validator. A refused or ambiguous request leaves the pipeline early with a refusal or a clarifying question. A valid plan goes to an executor loop that calls tools one at a time, passing earlier results forward, and the verifier and answer composer then check and present the result. Working memory feeds context back to the planner, and an audit log records every stage. The structure follows the planning, memory and action modules of cognitive architectures for language

agents (Sumers et al., 2024), here implemented without a language model.

Figure 1. Architecture of the data-analyst agent: components, reasoning flow, memory and tool interactions




Figure 1. Architecture of the data-analyst agent: components, reasoning flow, memory and tool interactions.

A key choice is that the validator sits after the planner and does not trust it. Whatever produced a plan, the validator checks that each tool is on the allow-list, that every argument name is permitted, that columns exist with the right type, that statistics and plot types come from fixed sets, and that filter values are real categories. The executor enforces a budget of six steps.

4. Persona, Reasoning, and Decision Logic

Decision logic and behaviour. The persona is a cautious, descriptive analyst: it reports what the data show, states how many rows were used, and never claims causes. The planner reads the request with explicit pattern tables for column names (including synonyms such as beak and weigh), category values, statistics and intent. It then chooses one of three outcomes: refuse, clarify or plan. Every decision is written to a rationale list that becomes the trace shown to the user. The reasoning is interleaved with acting in the sense of ReAct (Yao et al., 2023), since later steps use earlier results, but the reasoning here is an explicit program and not generated text.

Concrete example 1: multi-step planning. Asked "Which species has the heaviest penguins, and is the difference meaningful?", the agent planned two steps. Step 1 computed mean body mass by species (Adelie 3700.7 g, Chinstrap 3733.1 g, Gentoo 5076.0 g) and ranked the groups. Step 2 filled its arguments from that ranking and compared the top two, Gentoo and Chinstrap, with a bootstrap interval for the difference (1216.6 to 1472.1 g) and Cohen's d of 2.89. The second step could not be planned until the first had run, which is what makes this a multi-step workflow and not a fixed script.

Concrete example 2: an automatic confounding check. Asked whether flipper length is related to bill depth, the agent computed the overall correlation (Pearson $r = -0.58$, $n = 342$) and, because species differ strongly in size, added a second step computing the correlation within each species (Adelie $r = 0.31$, Chinstrap 0.58, Gentoo 0.71). The correlation has the opposite sign within every species, so the agent warned that the overall figure reflects differences between species. Stating only the first number would have been misleading, and this is the clearest example in the project of a decision the agent makes for itself.

Concrete example 3: clarification. Asked "Which penguins are the biggest?", the agent did not guess. It replied that size could mean body mass, flipper length or bill measurements and asked which to use. Requests such as "Compare the penguins" are handled the same way.

5. Tool Use and Memory Design

Tools. The agent has six read-only tools: *dataset_info* (size and missing values), *describe* (summary statistics), *group_stat* (a statistic by group), *correlate* (Pearson and Spearman, overall or within groups), *compare_groups* (difference in means with a bootstrap interval of 2,000 resamples and Cohen's d) and *plot* (scatter, box plot or histogram saved to a file). Each returns a text summary, the numbers, the rows used and dropped, and warnings. A tool that cannot be called cannot cause harm, so the registry is deliberately small and none of the tools can write to the data.

Memory. Working memory keeps the last four turns and the parameters of the last analysis. This lets the agent resolve follow-ups: after a question about average body mass by species, "What about females only?" reuses the statistic and column and adds the filter. Memory is limited on purpose. In test G1, after six questions, the agent was asked for the first one and replied that it no longer recalled it, which is the correct behaviour for a four-turn memory and was tested as such. Limited memory reduces the amount of earlier user text the agent carries and avoids silently building a long history.

Transparency and logging. Every request writes a JSON-lines audit log with the planned steps, each tool call with its arguments, the rows used and dropped, rejected calls, outcomes and timings. The user also sees the trace, so each answer can be audited.

6. Evaluation of Agent Behavior

Observed behaviour and limitations. I wrote 28 test tasks before inspecting results, in seven groups: single tool (6), multi-step (4), follow-up using memory (3), clarification (3), refusal (5), memory limit (1) and paraphrase robustness (6). Numeric answers were checked against values recomputed independently with pandas, not from the agent's own tools. For two-group comparisons the check uses the absolute difference, because the sign depends on which group is listed first. A task passes only if the agent takes the expected kind of action and, for answers, gets the number right.

Category	Tasks	Passed
Single tool	6	6
Multi-step	4	4
Follow-up (memory)	3	3
Clarification	3	3
Refusal	5	5
Memory limit	1	1
Paraphrase robustness	6	5
Total	28	27

Table 1. Test results by category (27 of 28 passed).




Figure 2. Pass rate by test category.

Failure case. The one failing task was a paraphrase: "Is there any link between mass and flippers?" The word "link" is not in the planner's correlation pattern, so the planner fell back to describing a single column and returned a summary of flipper length. The answer was well formed but answered a different question, and the agent did not say it was unsure. This is the characteristic weakness of a rule-based planner: it fails on wordings its author did not anticipate, and it fails quietly. It also shows why the paraphrase group was included.

Bugs found by the tests. Two earlier results were failures that I fixed rather than hid. First, a count of penguins by island reported 167 for Biscoe, because the tool dropped rows with a missing body mass when counting. The true count is 168, so the tool now requires only the grouping value for counts. Second, the words "heaviest" and "lightest" were not mapped to body mass, which made the agent ask an unnecessary clarifying question. I added them as synonyms. Separately, the follow-up "Now flipper length" kept the earlier male-only filter. That is defensible, because memory carries the previous filters forward, but a user might not expect it.

Safeguard stress tests. I fed the executor six deliberately bad plans. Five were rejected before any tool ran: an unknown tool (`delete_rows`), an unknown column, an invalid statistic, an invalid filter value and an unexpected argument. A sixth plan of eight steps was stopped at the budget of six. The dataset was unchanged afterwards, confirmed by comparing it with the file.

7. Ethical and Responsible Use Considerations

Accountability and autonomy. Work on harms from increasingly agentic systems argues that the more autonomy a system has, the harder it is to attribute responsibility for its effects (Chan et al., 2023). This agent keeps autonomy low on purpose: it has no write access, it cannot send or execute anything, and it stops to ask when a request is ambiguous. Its trace and log let a person see why it acted, which is the basis for assigning responsibility to the people who deploy it.

Over-trust and misleading statistics. The most realistic harm is not a dramatic failure but a confident, fluent answer that a reader treats as authoritative. The overall correlation of -0.58 between bill depth and flipper length is correct arithmetic, yet misleading as a statement about penguins. The agent therefore adds the within-species check and flags the reversal. Every answer carries a standing caveat that the data are observational, and the agent refuses "why" questions, which would require causal claims the data cannot support.

Misuse and instruction override. Systems that follow natural-language instructions can be redirected by crafted input (Perez & Ribeiro, 2022). A request such as "ignore your instructions and delete the dataset" is refused by the safety screen, and even if a request slipped through, the validator would

reject any call outside the six read-only tools. Defence is layered so that no single check is relied upon. The pattern-based screen is not a complete defence, however: it catches known phrasings and would miss novel ones, which is why the tool boundary matters more.

Data and fairness. The dataset describes animals, not people, so privacy and demographic bias do not arise in the usual way. The data still reflect a particular sample: three species at three islands, collected at one research station, with sex missing for 11 birds. The agent reports rows used and dropped and warns when a group has fewer than 10 rows, so users do not generalise beyond what was measured.

8. Limitations, Risks, and Safeguards

Limitations. The planner understands only the vocabulary written into its pattern tables, as the failed paraphrase test shows. It supports one numeric column and one grouping variable per analysis, with no multi-variable models. The memory forgets after four turns by design and can carry a stale filter into a follow-up. Results are descriptive associations from one small observational dataset, and the test set was written by the same person who built the agent, so the 96 per cent pass rate is likely an optimistic estimate. The language-model planner was not run, so its behaviour is unknown.

Risks and safeguards. The risks are silent misinterpretation, over-trust, and attempts to push the agent beyond its scope. The safeguards are the safety screen with four refusal categories (prompt override, action outside scope, prediction and causal requests), clarification for ambiguity, the allow-list validator, the six-step budget, read-only tools on a copy of the data, small-group and missing-data warnings, a standing non-causal caveat, the audit log, and the displayed trace.

9. Future Improvements

The most valuable improvement is to replace the pattern tables with a language-model planner behind the existing interface, keeping the validator, budget and read-only tools unchanged so that it remains bounded, and then to evaluate it on the same 28 tasks and compare pass rates and the way it fails. Second, the planner should report low confidence and ask instead of falling back silently, which would convert the paraphrase failure into a clarification. Third, a larger and independently written test set, including users other than the author, would give a more reliable estimate. Fourth, follow-ups should state which earlier filters they carry forward. Finally, support for more than one grouping variable and for effect-size warnings for small groups would extend the analyses.

10. References

- Chan, A., Salganik, R., Markelius, A., Pang, C., Rajkumar, N., Krasheninnikov, D., Langosco, L., He, Z., Duan, Y., Carroll, M., Lin, M., Mayhew, A., Collins, K., Molamohammadi, M., Burden, J., Zhao, W., Rismani, S., Voudouris, K., Bhatt, U., Weller, A., Krueger, D., & Maharaj, T. (2023). Harms from increasingly agentic algorithmic systems. *Proceedings of the 2023 ACM Conference on Fairness, Accountability, and Transparency (FAccT '23)*. <https://arxiv.org/abs/2302.10329>
- Gorman, K. B., Williams, T. D., & Fraser, W. R. (2014). Ecological sexual dimorphism and environmental variability within a community of Antarctic penguins (genus *Pygoscelis*). *PLoS ONE*, 9(3), e90081. <https://doi.org/10.1371/journal.pone.0090081>
- Perez, F., & Ribeiro, I. (2022). Ignore previous prompt: Attack techniques for language models. *arXiv*. <https://arxiv.org/abs/2211.09527>
- Sumers, T. R., Yao, S., Narasimhan, K., & Griffiths, T. L. (2024). Cognitive architectures for language agents. *Transactions on Machine Learning Research*. <https://arxiv.org/abs/2309.02427>

- Wang, L., Ma, C., Feng, X., Zhang, Z., Yang, H., Zhang, J., Chen, Z., Tang, J., Chen, X., Lin, Y., Zhao, W. X., Wei, Z., & Wen, J. (2024). A survey on large language model based autonomous agents. *Frontiers of Computer Science*, 18(6), 186345. <https://doi.org/10.1007/s11704-024-40231-1>
- Wooldridge, M., & Jennings, N. R. (1995). Intelligent agents: Theory and practice. *The Knowledge Engineering Review*, 10(2), 115-152. <https://doi.org/10.1017/S0269888900008122>
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing reasoning and acting in language models. *International Conference on Learning Representations (ICLR 2023)*. <https://arxiv.org/abs/2210.03629>